

# Page Table Structure

- Hierarchical Paging
- Hashed Page Tables
- Inverted Page Tables

# Hierarchical Page Tables

- **Problem:** for very large logical address spaces ( $2^{32}$  and  $2^{64}$ ) the page table itself becomes very large.
- **Solution:** break up the logical address space into multiple page tables.
- A simple technique is a two-level page table.


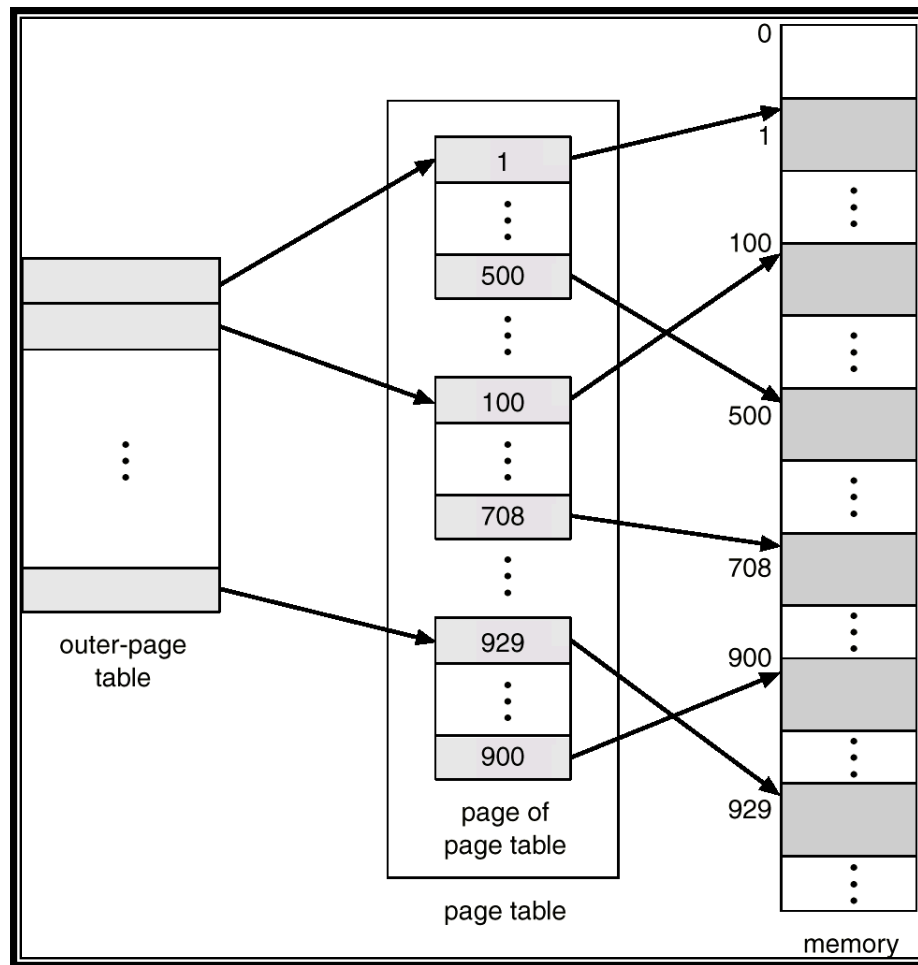
# Two-Level Paging Example

- A logical address (on 32-bit machine with 4K page size) is divided into:
  - ❑ a page number consisting of 20 bits.
  - ❑ a page offset consisting of 12 bits.
- Since the page table is paged, the page number is further divided into:
  - ❑ a 10-bit page number.
  - ❑ a 10-bit page offset.
- Thus, a logical address is as follows:

page number		page offset
$p_1$	$p_2$	$d$
10	10	12

where  $p_1$  is an index into the **outer page table**, and  $p_2$  is the index into the **inner page table**.


# Two-Level Page-Table Scheme





**PTBR**: base address of the **outer-page table**  
**PTLR**: number of the **inner-page tables**

# Address-Translation Scheme

- Address-translation scheme for a two-level 32-bit paging architecture




# Three-level Paging Scheme for 64-bit address space



# Hashed Page Table Scheme

- Common in address spaces  $> 32$  bits.
- The logical page number is hashed into a page table.
  - ❑ This page table contains a **chain of elements** in the form of  **$\langle \text{page\#}, \text{frame\#} \rangle$**  hashing to the same location.
- Logical page number is compared in this chain searching for a match.
  - ❑ If a match is found, the corresponding physical frame# is extracted.

# Hashed Page Table





# Hash tables

## Linear Probing

0	
1	141
2	621
3	123
4	
5	465
6	46
7	67
8	88
9	288
10	390
11	91
12	152
13	572
14	734
15	155
16	
17	
18	178
19	399

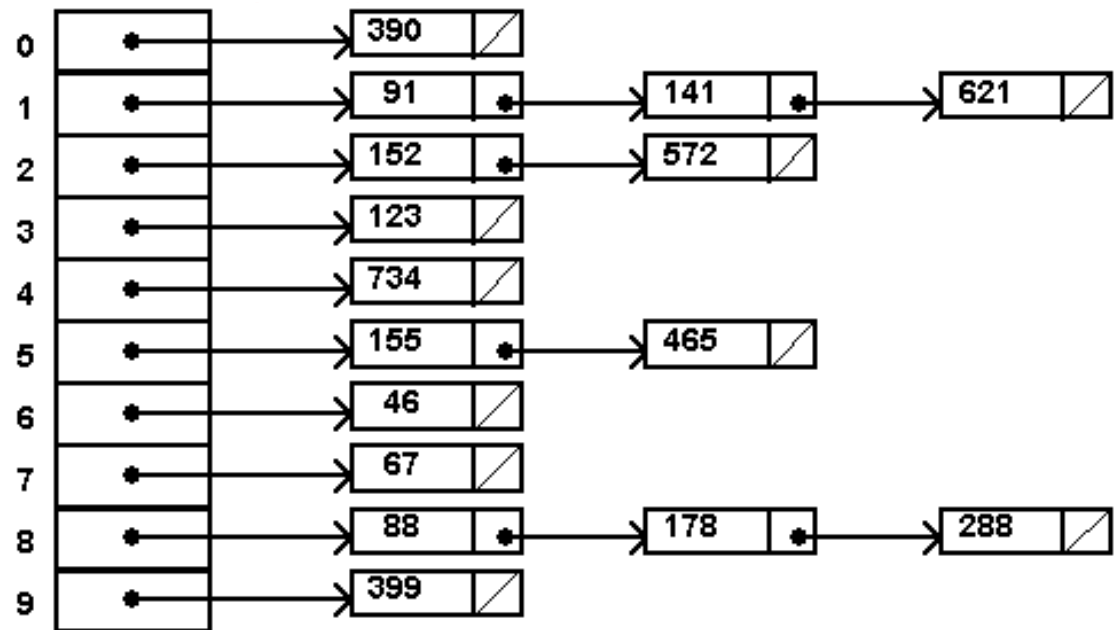
$67 \% 20 = 7$   
 $46 \% 20 = 6$   
 $88 \% 20 = 8$   
 $91 \% 20 = 11$

...

$734 \% 20 = 14$

Hash Function

## Chained Overflow



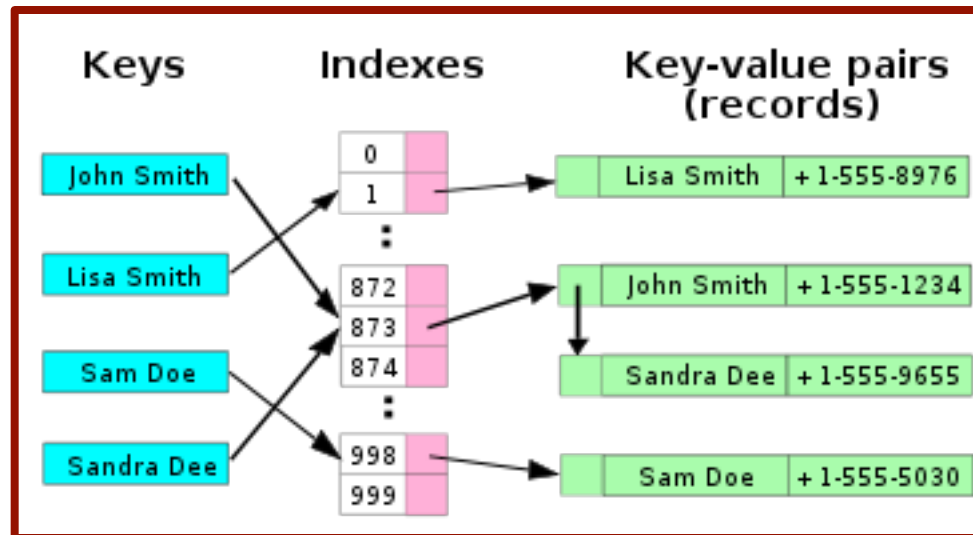
$67 \% 10 = 7$   
 $46 \% 10 = 6$   
 $88 \% 10 = 8$   
 $91 \% 10 = 1$

...  
 $734 \% 10 = 4$

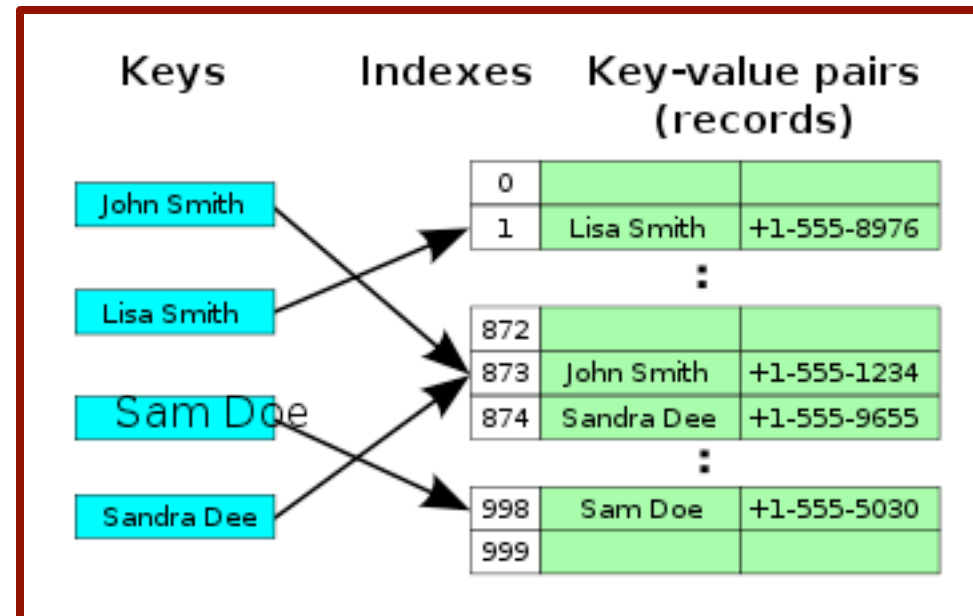
**Only Key,  
Value not shown.  
(Key, Value) pairs.**

# Hash Table Example of Telephone Directory

Hash collision resolved by chaining



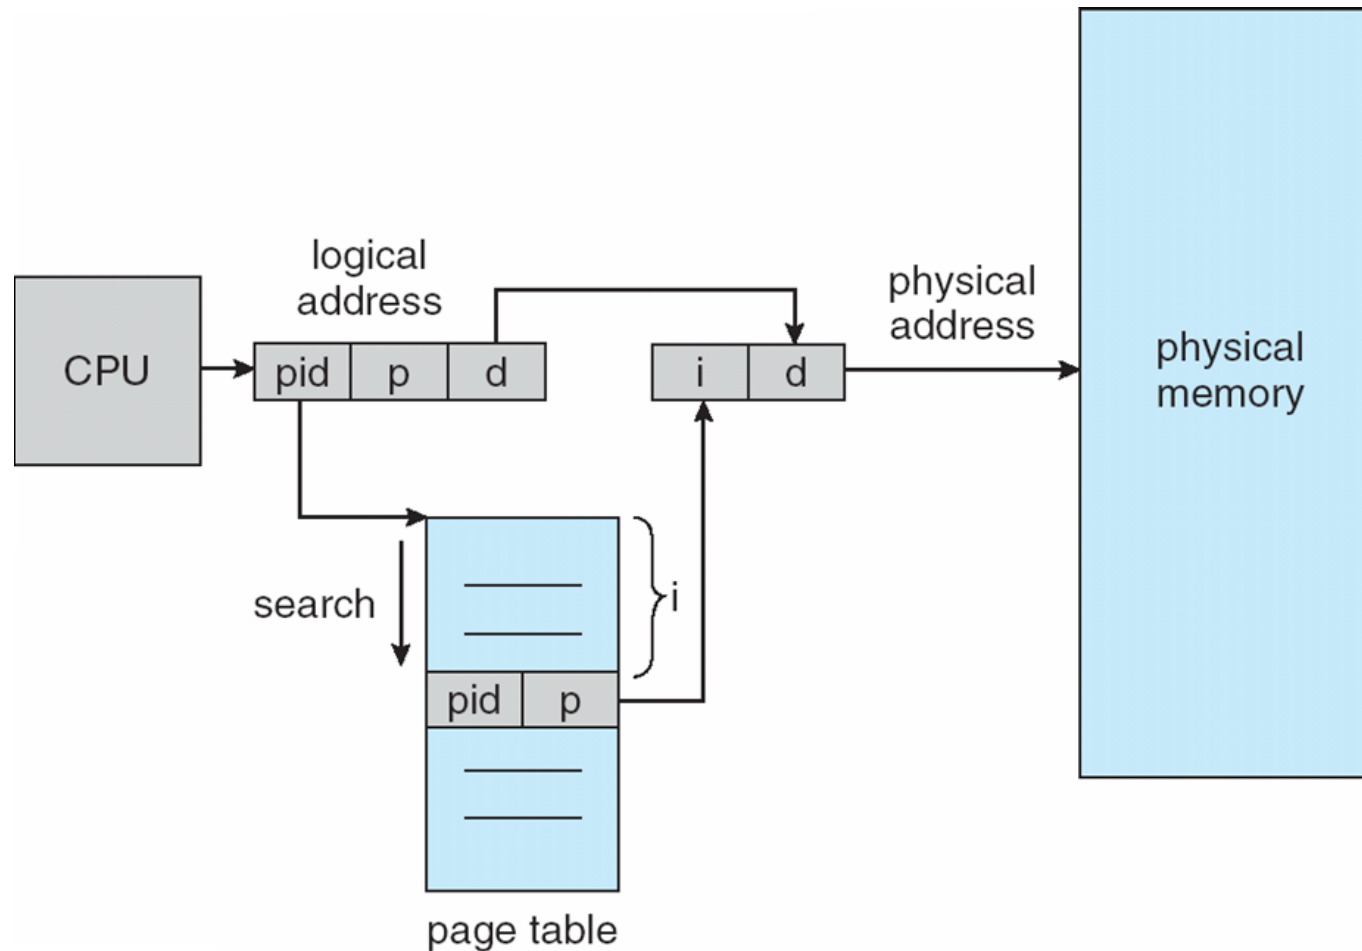
Hash collision resolved by linear probing (interval = 1)



# Inverted Page Table

- One entry for each frame of physical memory.
- Entry consists of the logical address of the page stored in that real memory location, with information about the process that owns that page.
- Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs.
- Use hash table to limit the search to one — or at most a few — page-table entries.

# Inverted Page Table Architecture



# Shared Pages

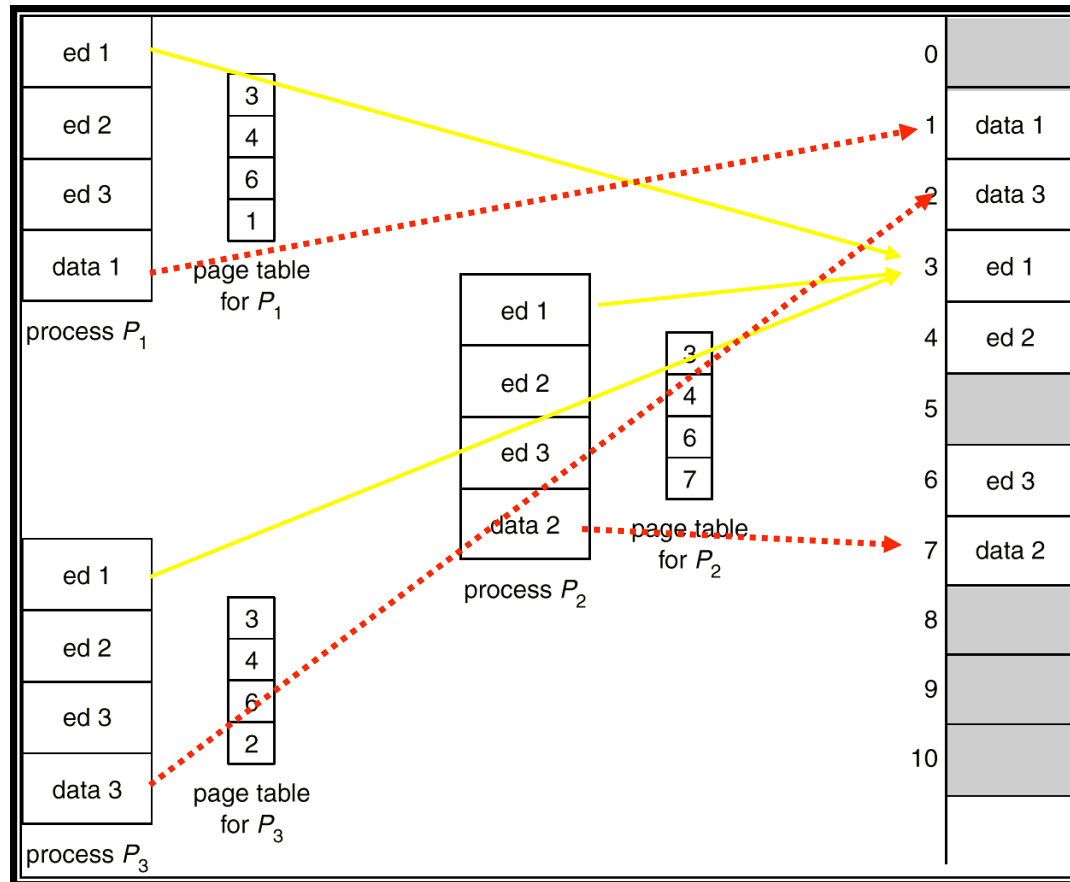
## ➤ Shared code

- ❑ One copy of **read-only (reentrant) code** shared among processes (i.e., text editors, compilers, window systems).
- ❑ Shared code must appear in same location in the logical address space of all processes

## ➤ Private code and data

- ❑ Each process keeps a separate copy of the code and data.
- ❑ The pages for the private code and data can appear anywhere in the logical address space.

# Shared Pages Example

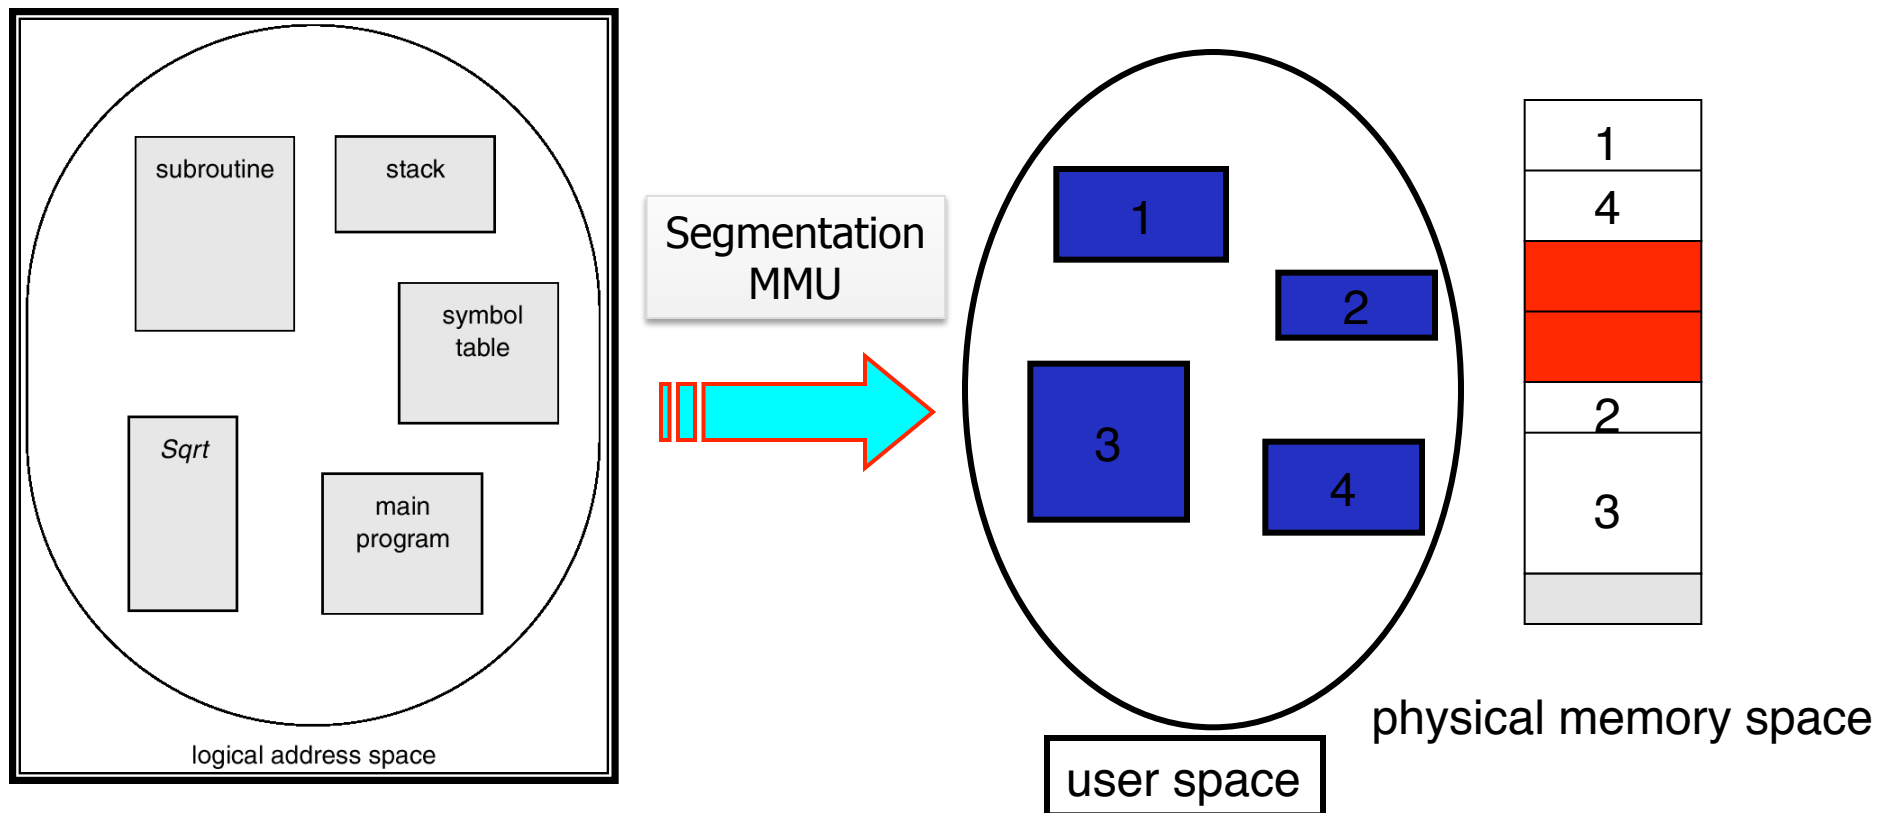


# Segmentation

- Memory-management scheme that supports user view of program.
- User does not look at program as an array of memory data
- A program is a collection of segments.
  - ❑ A segment is a logical unit such as:
    - main program,
    - procedure,
    - function,
    - method,
    - object,
    - local variables, global variables,
    - common block,
    - stack,
    - symbol table,
    - arrays

# User's View of a Program

## & Logical View of Segmentation





# Segmentation Architecture

- Logical address consists of a two tuple:  
 $\langle \text{segment-number}, \text{offset} \rangle,$
- *Segment table* – maps two-dimensional logical addresses; each table entry has:
  - ❑ *base* – contains the starting physical address where the segments reside in memory.
  - ❑ *limit* – specifies the length of the segment.
- *Segment-table base register (STBR)* points to the segment table's location in memory.
- *Segment-table length register (STLR)* indicates number of segments used by a program;  
segment number  $s$  is legal if  $s < \text{STLR}$ .

# Segmentation Architecture (Cont.)

## Features:

### ➤ Relocation.

- ❑ dynamic
- ❑ by segment table

### ➤ Sharing.

- ❑ shared segments
- ❑ same segment number

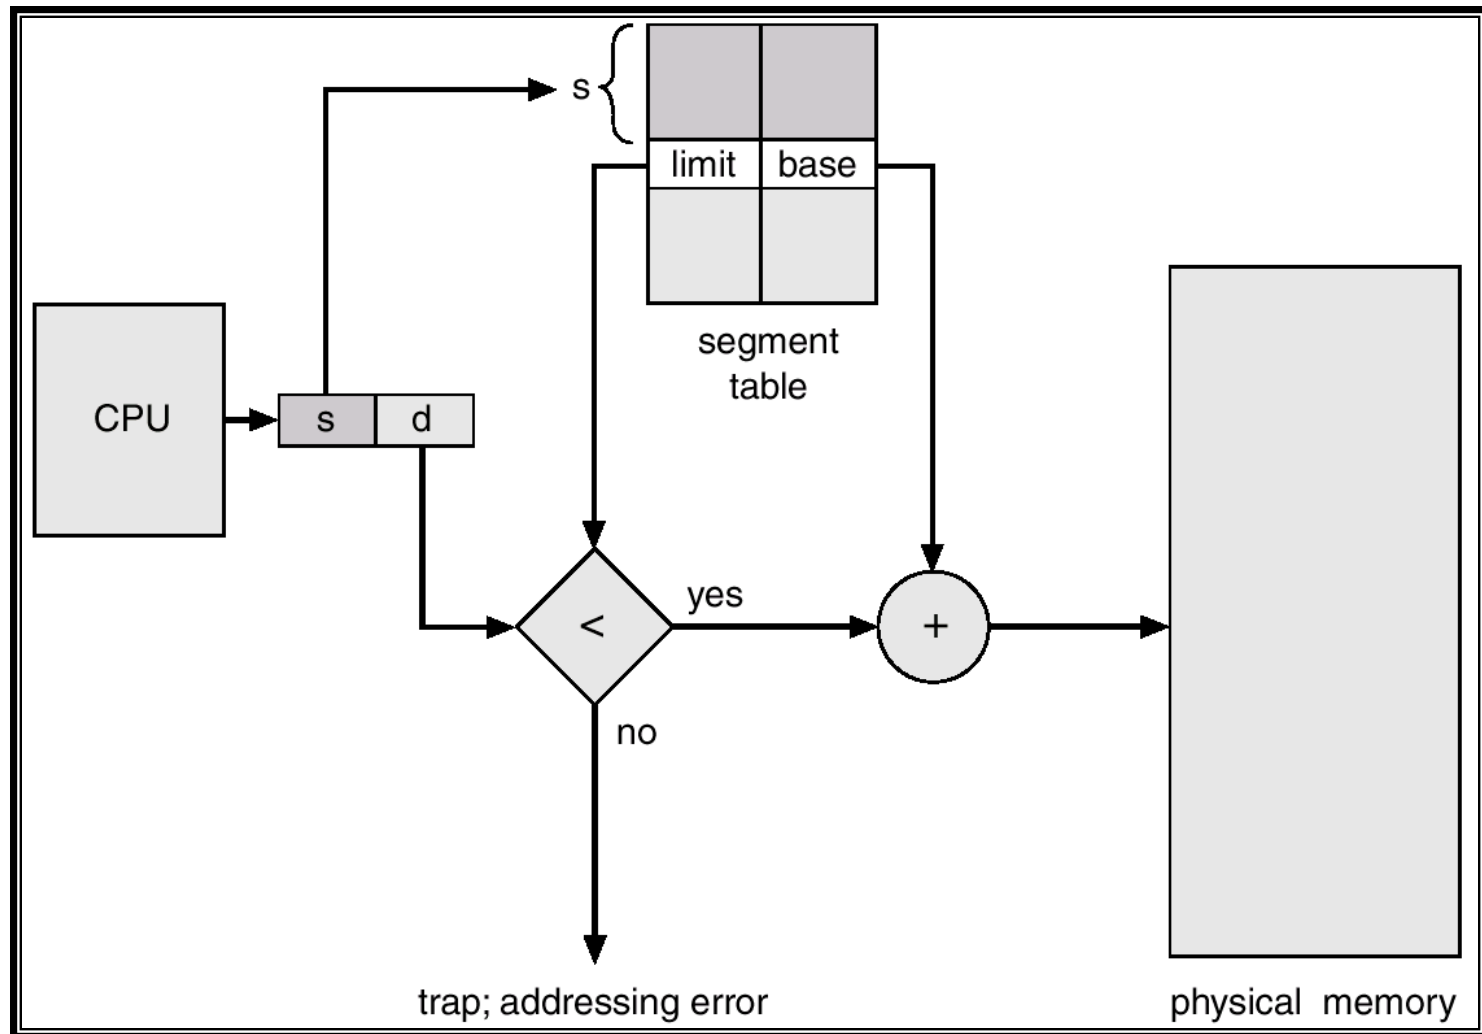
### ➤ Allocation.

- ❑ first fit/best fit
- ❑ external fragmentation

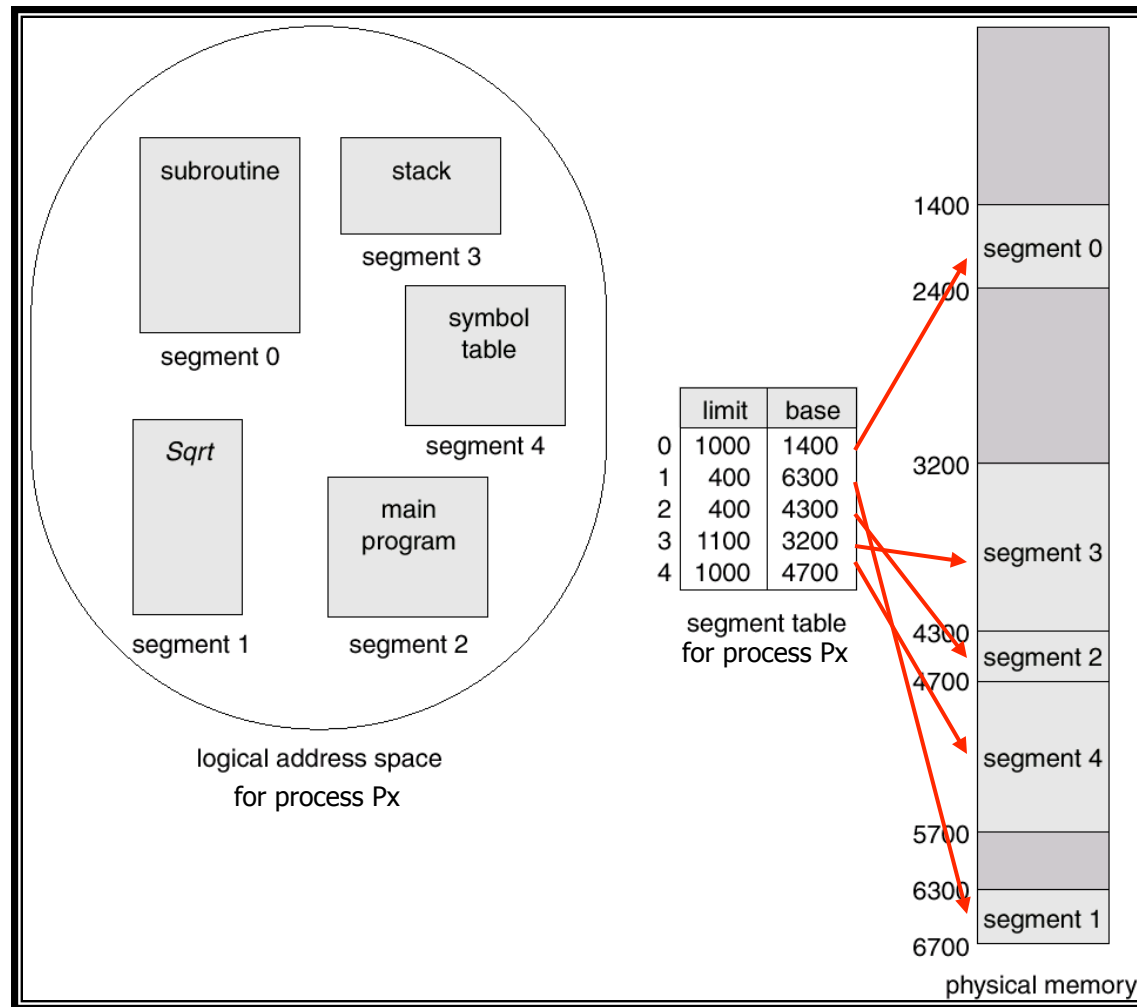
# Segmentation Architecture (Cont.)

- **Protection.** With each entry in segment table associate:
  - ❑ validation bit = 0  $\Rightarrow$  illegal segment
  - ❑ read/write/execute privileges
- Protection bits associated with segments; code sharing occurs at segment level.
- Since segments vary in length, memory allocation is a dynamic storage-allocation problem.
- A segmentation example is shown in the following diagram

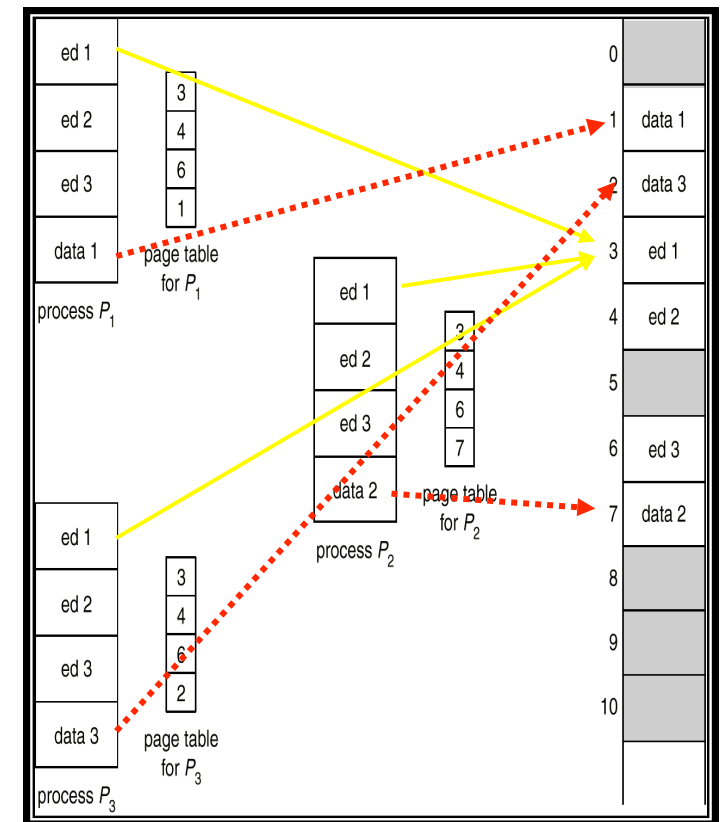
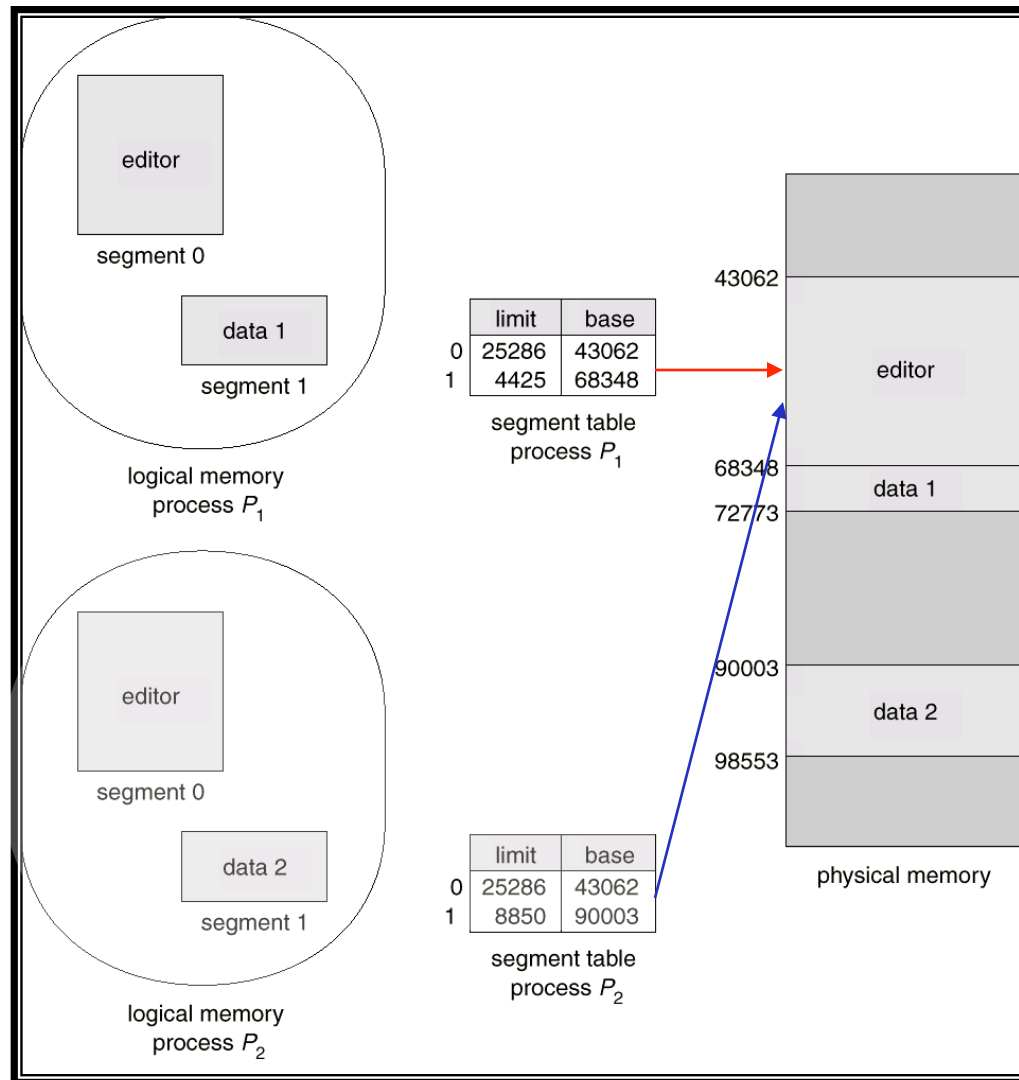
# Segmentation Hardware



# Example of Segmentation



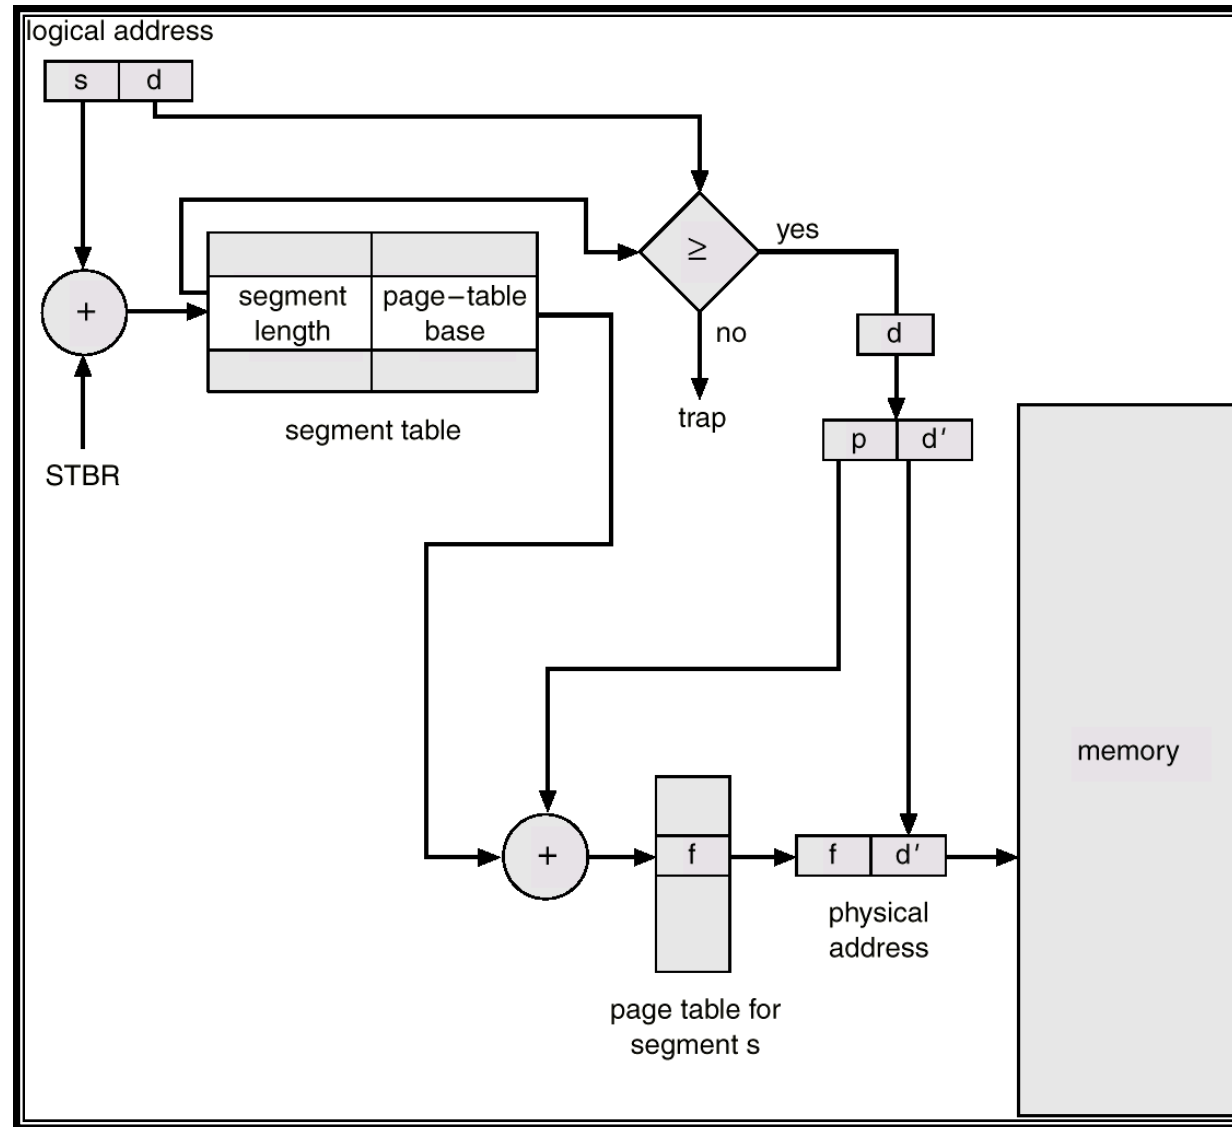
# Sharing of Segments



# Segmentation with Paging

- The MULTICS and Intel 80386 systems solved problems of external fragmentation and lengthy search time for empty slot by **paging the segments**.
- Solution differs from pure segmentation in that the **segment-table** entry contains not the base address of the segment, but rather the **base address of a *page table* for this segment**.

# MULTICS Address Translation Scheme





# Intel 80386

- The Intel 386 uses **segmentation with paging** for memory management with a **two-level paging** scheme.
- Logical address consists of a **segment selector** (in one of the **segment registers** of CPU) and **32 bits address** that is generated on the address bus.
- Each process can have up to **8000 private segments** and up to **8000 shared segments** to be shared by other processes.
- Each segment can have up to **1 million pages** (each **page is 4K bytes**) hence the size of a segment can be up to **4 Giga bytes**.
- The **pages** of a segment **are scattered** throughout the memory frames.
- A 32 bit address points to the start of a segment of pages and a 32 bit address from address bus is used as the offset in that segment. This offset is used to access the locations in the pages of the segment.
- Intel 386 has hardware to support virtual memory management effectively.

# Segmentation with Paging - Intel 386

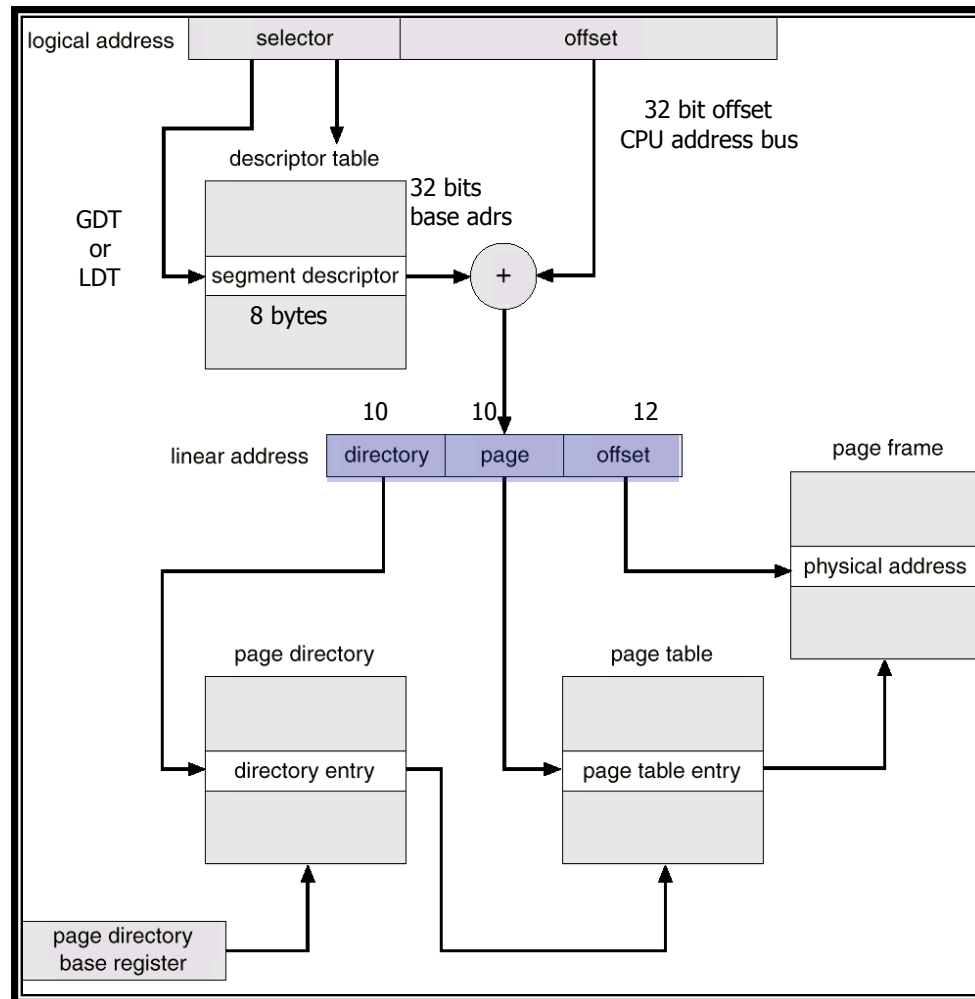
- The Intel 386 uses segmentation with paging for memory management with a two-level paging scheme.

## Intel 80386 address translation

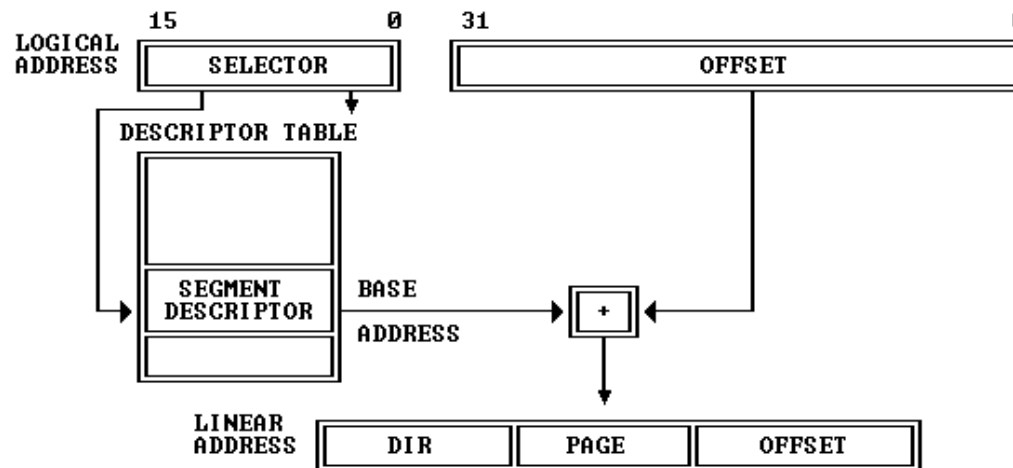


Selector part of logical address.


Stored in one of the CPU's  
segment registers:  
CS, DS, SS, ES, ..




# 80386 Architecture for MMU



## DESCRIPTORS USED FOR APPLICATIONS CODE AND DATA SEGMENTS



## DESCRIPTORS USED FOR SPECIAL SYSTEM SEGMENTS



- A - ACCESSED
- AUL - AVAILABLE FOR USE BY SYSTEMS PROGRAMMERS
- DPL - DESCRIPTOR PRIVILEGE LEVEL
- G - GRANULARITY
- P - SEGMENT PRESENT

# 80386 Architecture


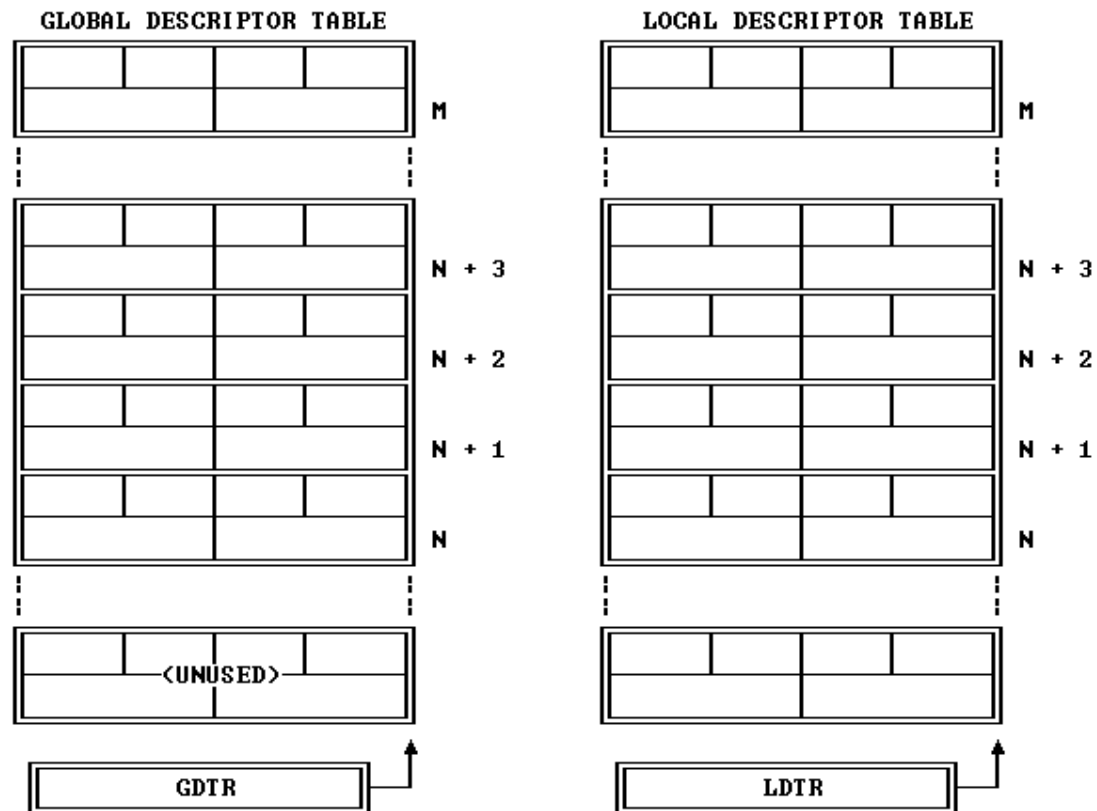




Figure 5-5. Descriptor Tables



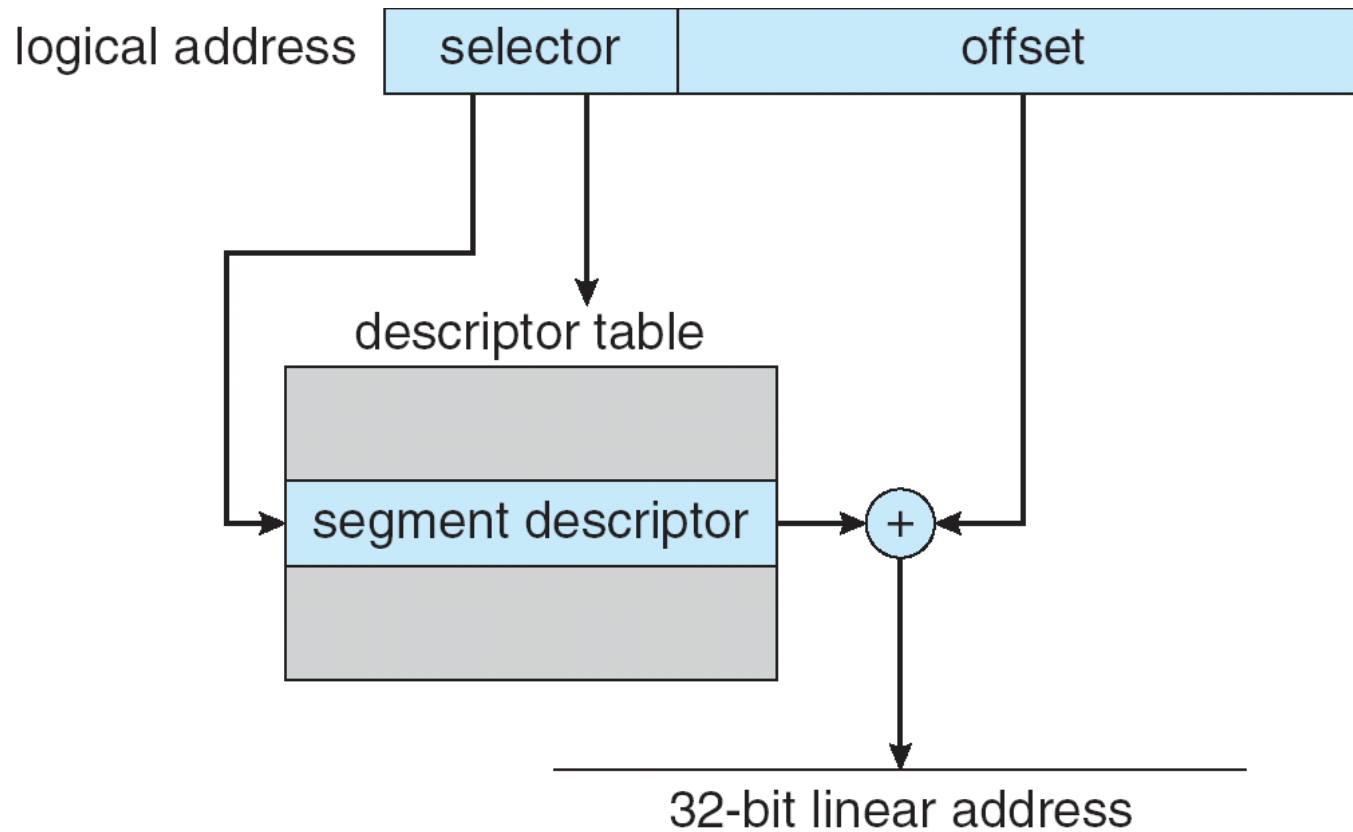
# Example: The Intel Pentium

- Supports both segmentation and segmentation with paging
- CPU generates logical address
  - Given to segmentation unit
    - ❖ Which produces linear addresses
  - Linear address given to paging unit
    - ❖ Which generates physical address in main memory
    - ❖ Paging units form equivalent of MMU

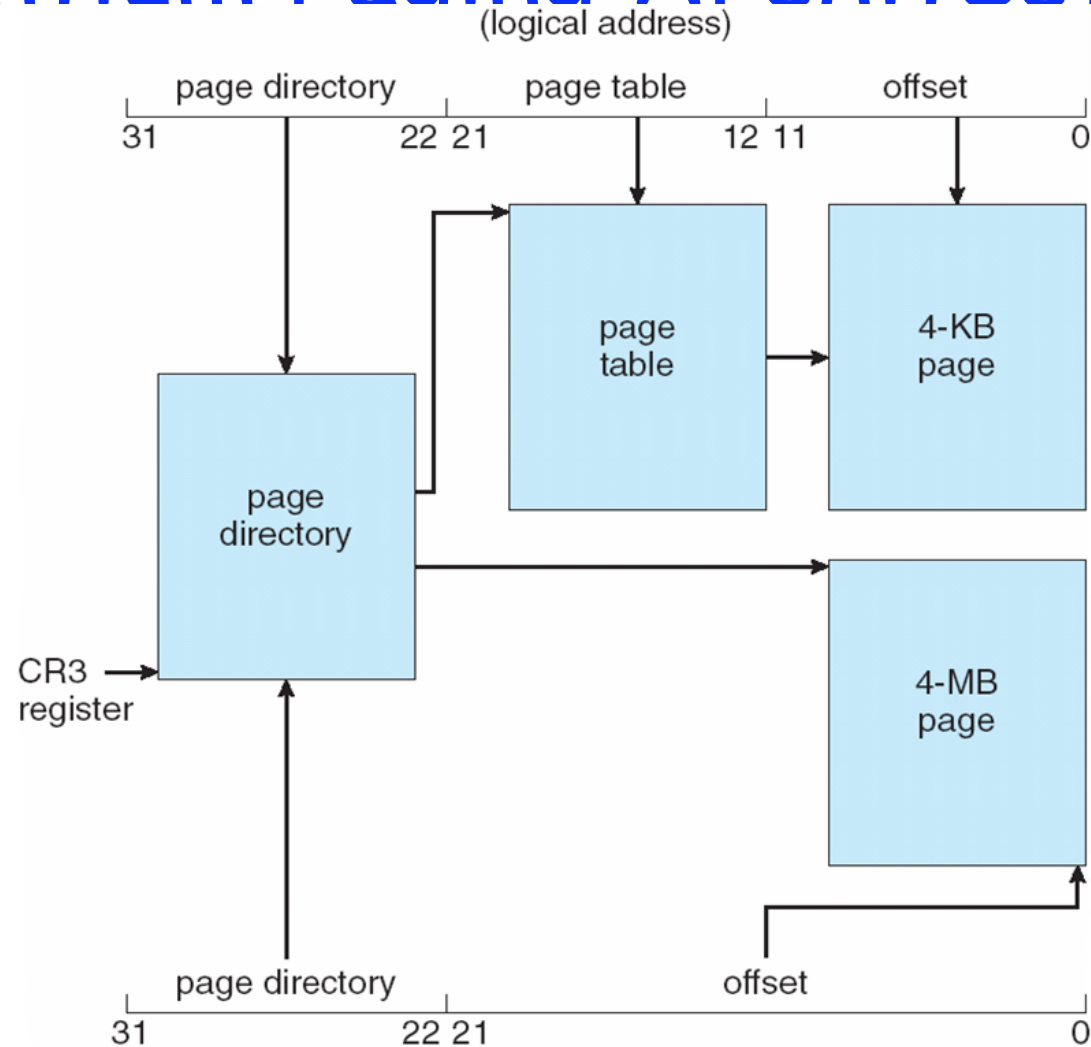
# Logical to Physical Address Translation in Pentium



# Intel Pentium Segmentation



# Pentium Paaina Architecture





# Linear Address in Linux

Broken into four parts:

global directory	middle directory	page table	offset
---------------------	---------------------	---------------	--------

# Three-level Paging in Linux

(linear address)

